

# CS250: An Elliptic Curve Cryptography Engine

Palmer Dabbelt and Kevin Linger

December 18, 2013

For our CS250 project, we implemented an elliptic curve cryptography engine that attaches via the Rocket coprocessor interface. Our accelerator was targeted at supporting the elliptic curve digital signature algorithm (ECDSA), a NIST standard widely used on the Internet for digital signatures. We built off of previous work by doing a more involved design-space exploration, focusing on hardware/software co-tuning. We were succesful in implementing the primary operation needed in ECDSA which is an elliptic curve point multiplication. A simulated version of the operation works with the Rocket Core, but the full toolchain did not run. With some algorithmic improvements to our accelerator, we can get tremendous improvements over a software implementation.

## 1 Elliptic Curve Cryptography

Digital signature algorithms[10] are a widely used class of algorithms that are central to the functioning of the modern Internet. DSA[7] (the original digital signature algorithm as standardized by NIST) is a specialization of the ElGamal signature scheme[3] that uses discrete logarithms over finite integer fields. Finite integer fields are particularly easy to compute with electronics, so this gives DSA the advantage of being an efficient signature algorithm.

The problem is that the discrete logarithm problem is not particularly strong over finite integer fields (specifically, there are known sub-exponential algorithms[1]). As shown in Table 1, DSA does not scale well to higher security levels. This is important because code breaking algorithms are getting more efficient, leading to the need for higher security levels, therefore imposing higher computational costs on every node in the system.

The commonly accepted solution to this problem is to change the field from integers to elliptic curves, which (as shown in Table 1) scale much better to higher security levels. While elliptic curves are significantly more difficult to build electronics for, this

Security Level	DSA	ECDSA	Ratio
80	1024	160	3:1
112	2048	224	6:1
128	3072	256	10:1
192	7680	384	32:1
256	15360	521	64:1

Table 1: Comparison of DSA and ECDSA[2]

will eventually be offset by sheer key size alone. The currently accepted security level is 128 bits, at which ECDSA already has a 10:1 advantage in key size. Shortening the key size becomes even more important, as we approach a world in which ubiquitous embedded systems must be able to securely transmit and receive data.

In addition to providing for faster cryptographic operations, this shortened key size plays an important role specific to the swarm related to key generation. It is already known[4] that we can't find sufficient entropy to generate RSA-style keys for the machines we have, much less for a network the size of that envisioned by the swarm project. The main problem with generating RSA keys is finding enough entropy to search for large prime numbers, a problem that goes away with ECDSA because the numbers are smaller and there's no necessity for them to be prime (there is an acceptability criteria, but the chances of a random number not meeting it are exceedingly low).

## 2 The Swarm

One particularly interesting application of public-key cryptography arose as part of the SwarmOS project, which was the original motivation for this accelerator. The SwarmOS is designed to be an operating system for the swarm, which consists of a large number ( $10^{10}$ ) of low-powered ( $\mu$ W) devices as well as a smaller number of more powerful servers.

One component of the SwarmOS is the Universal Dataplane, which is designed to serve as a combined communication and storage mechanism for swarm devices. The general idea is that data will be put into the dataplane in one place by sensors, and can then be removed anywhere else in the world by users. This involves a significant amount of data being between untrusted servers, and as such the stance we decided to take was to sign and encrypt all data all the time.

Unfortunately, ECDSA requires a significant amount of energy to perform each signature operation (encryption is less scary because you can cheat a bit like PGP does) – software ECDSA is mJ per signature, which leads to very low bandwidth on a  $\mu\text{W}$  machine. The only way to get the required  $1000\times$  energy per operation improvement is to implement ECDSA in hardware, which is essentially what this project is.

### 3 Elliptic Curve Arithmetic

The compute-intensive operation in ECDSA is an elliptic curve (EC) point multiplication. An EC multiply is a series of EC doubles and additions, which are much different than the respective integer operations. These operations are defined geometrically; e.g. to double a point on an EC, a tangent line at that point is drawn which will intersect the curve at exactly one other point. The result of the double operation is that intersection point reflected over the X-axis. An example of the geometric definition of point addition is shown in Figure 1, obtained from Certicom, which involves drawing a line between two points on the curve, tracing its intersection with the curve, and then a reflection over the X axis – point doubling is almost exactly the same, except with a tangent instead of a line between the two points (you can think of the tangent like the limit of a line between two points as those two points approach each other).

The point add operation is described by the equation shown in Figure 2, and the point double operation is described by the equation shown in Figure 3. In these equations all operations are modulo the order of the curve, which is the same order of magnitude of the prime modulus of the curve (but it's a different number). Additionally, the curve constant  $a$  is as defined by NIST.

Each of these operations also has a modular operation at the end, where the modulus is a large prime number that represents the order of the curve. This




Figure 1: Graphical definition of point add

$$\lambda = \frac{P_y - Q_y}{P_x - Q_x}$$

$$R_x = \lambda^2 - P_x - Q_x$$

$$R_y = \lambda(P_x - R_x) - P_y$$

Figure 2: Point addition formula

makes multiplication and division much more difficult: in an accelerator where an addition takes two cycles, a modular multiply will take on the order of a thousand cycles[6]. A division is now a modular inverse followed by a modular multiplication.

It is possible to formulate the discrete logarithmic problem over elliptic curves by repeating addition as described above. Under standard hardness assumptions, this formulation of the discrete logarithm problem is asymmetric (actually, this is one of the standard hardness assumptions). The general idea is simple: point multiplication is defined as repeated addition. Without point doubling, for an  $N$ -bit number we would have to repeat addition  $O(2^N)$  times in order to do a multiplication. Having a formula for point doubling allows us to reduce this to  $O(N)$  point doubles along with  $O(N)$  point adds by storing some intermediate calculations along the way, just like the long multiplication algorithm taught in grade school. There is no point-halving operation, so performing the inverse of a point multiplication still requires  $O(2^N)$  operations – that's where the asymmetry comes from.

$$\lambda = \frac{3 * P_x^2 + a}{2P_y}$$

$$R_x = \lambda^2 - P_x - Q_x$$

$$R_y = \lambda(P_x - R_x) - P_y$$

Figure 3: Point double formula

The full ECDSA algorithm essentially consists of doing a point multiplication to generate the private and public keys – the private key is a random integer and the public key is that integer multiplied by a well-known point, known as the generator point. The remainder of ECDSA consists of a few mathematical operations that generate enough information to verify a signature (shown in Figure 4), and then the inverse of that operation to verify a signature (shown in Figure 5). In these figures, “\*” is a point multiplication, while “×” is a modular multiplication – there is about a factor of 1000 speed difference between the two.

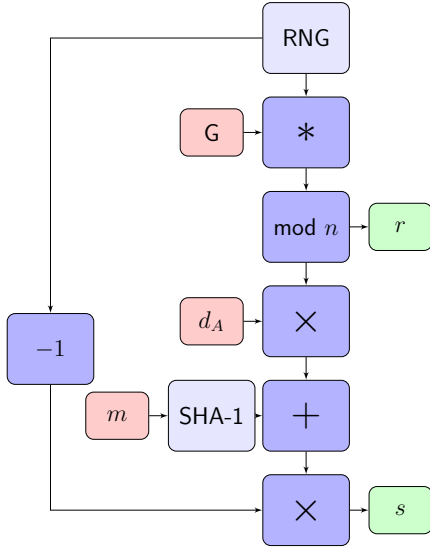


Figure 4: ECDSA Signature Algorithm

Due to the large difference in complexity the point multiplication algorithm and the remainder of the operations, the main focus of building a fast ECDSA implementation is building a fast point multiplier

The EC point multiplication is essentially a hierarchy of control logic that has the integer arithmetic at the base. This was the motivation of the hardware/software cotuning: at what level of the hierarchy should the hardware/software division be made?




Figure 5: ECDSA Verification Algorithm

## 4 Implementing ECDSA

Unlike finite integer fields, elliptic curves do not map directly to the hardware present in current microprocessors[6]. This suggests that a hardware accelerator designed to compute over elliptic curves should lead to a significantly more efficient implementation than in software alone. ECDSA requires more computation per bit than DSA, so a hardware accelerator is necessary to fully take advantage of the smaller key size.

Previous projects[9] have shown that ECC hardware can be implemented efficiently, but were limited in the amount of design-space exploration that was attempted. [8] explored the tradeoffs of different algorithms for computation, but did not explore the design space of any particular algorithm, which seems to be the extent of the design space exploration performed. In addition, the current state of the art appears to consist solely of FPGA implementations. Unfortunately these studies don't present power numbers, which is problematic because energy is the main limiting factor for swarm applications.

Figure 4 shows a flowchart for the signing operation and Figure 5 shows a flowchart for the verification operation. Both of these operations are dominated (approximately 99.9% of the runtime) by the point multiplication blocks. As such, most of the effort of this project was spent building a point multiplication block. That said, point multiplication involves repeating point doubling and addition, which themselves involve repeating modular multiplication and inversion.

## 5 Software Test Bench

The primary goal of this project was to explore the tradeoffs involved in adding different amounts of hardware support for ECDSA to a pre-existing chip. In order to achieve this, we needed to have the ability to generate a myriad of different software configurations that correspond to each of the accelerator instances we generated.

The current industry standard cryptographic software, OpenSSL, is very difficult to hack on. Instead of messing around with OpenSSL, we decided to instead write our own implementation of ECDSA with the explicit goal of being parameterizable. We could then generate an instance of this software that corresponds with every accelerator configuration that made sense and fully test our hardware.

Unfortunately, this software implementation proved to be the weakest point of this entire project. As shown in Table 2, our code is hundreds of times slower than OpenSSL – this is after a significant amount of tuning, our first implementations were hundreds of thousands of times slower than OpenSSL! This trouble can be attributed to two shortcomings: algorithms and tuning.

OpenSSL uses a number of significantly more complicated cryptographic operations than our code does. Specifically, it uses a projective point representation that allows for the conversion of a large number of modular inversions to a single inversion at the end of the multiplication. Implementing these algorithmic improvements should show significant speedups for both our software and hardware, so it’s not entirely a bad thing. We estimate about a  $20\times$  gain can be had from algorithmic improvements alone!

In addition to using better algorithms, OpenSSL is implemented better than our software was. Specifically, OpenSSL contains hand-tuned assembly for a number of critical paths that compilers appear to generate bad code for (specifically thinking of carry

chains for big integer operations). Part of the problem that our implementation had was probably related to the fact that it attempts to be parameterizable, which almost certainly defeats a lot of compiler heuristics (It would be great to tell a compiler to generate 4 different optimized versions of a function, sweeping across the 4 different values that one particularly interesting input can have).

Despite these shortcomings, we don’t believe that the software troubles completely invalidate the results from this project. We spent about the same amount of time on both the software and hardware. More importantly, we have numbers comparing against the full-strength OpenSSL code running on x86 (a platform that OpenSSL is highly optimized for[5]). These numbers are probably the most reasonable to compare our ASIC implementation against, as they’re the most realistic picture of the current state of the art.

## 6 ASIC Implementation

The ASIC implementation used the same algorithms as the software implementation, and as such looks very similar. Most ASICs end up being limited by memory bandwidth, but ECDSA is so computationally expensive we’re not even close to being bandwidth limited (we move roughly one bit per thousand cycles). This has the very nice property of allowing us to completely ignore memory, which means that the accelerator RTL ends up looking shockingly similar to the software implementation.

Our coprocessor interface is very simple – essentially we relied on the fact that almost no data was moved to avoid building anything complicated. A block diagram of the interface is shown in Figure 6. The general idea is that the coprocessor has two source registers, A and B and one destination register, C.

The host writes the data it would like to have computed to the source registers by emitting a copy instruction a number of times (depending on the size of the data type). The host then emits an assembly instruction that signifies what sort of operation it would like done. The accelerator’s control logic then sequences the functional units in the appropriate order, eventually writing back the result to the destination register. The host can then emit assembly instructions to read the results back – if the host emits a read operation too early it will block until the results are ready.

While this coprocessor interface has very high over-

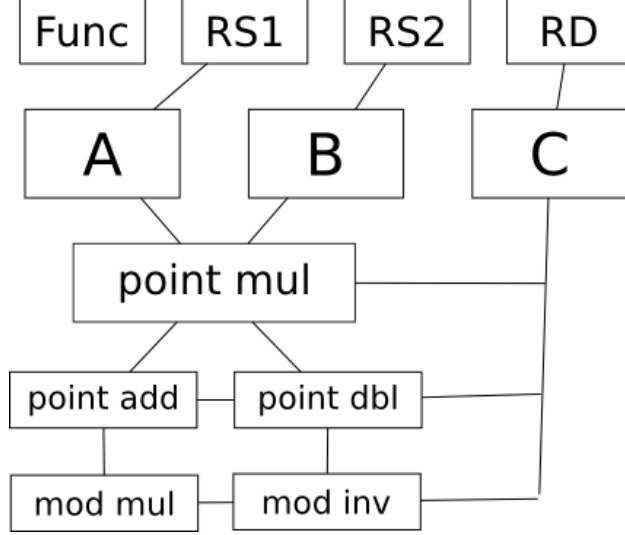


Figure 6: Rocket Coprocessor Interface

head, it turns out that our software implementation was sufficiently bad that doing any fine-grained accelerator operations ended up being pointless because of the software overhead (in addition to the interface overhead). For the interesting configurations the interface was very low overhead so it didn't end up causing too much trouble.

Our ASIC implementation of ECDSA performs all elliptic field arithmetic, including point multiplication, addition, doubling, as well as modular multiplication and modular inversion. Our hardware supports key sizes up to 256 bits, as that is the suggested security level.

We created instructions to expose each of the operations listed above. Operands are pushed 64-bits at a time. We opted to have operands pushed over the two 64-bit buses instead of having the accelerator fetch operands from memory. This primarily comes at the cost of register space in the rocket core, as it only takes four to twelve instructions to push all operands to the accelerator. In addition to setting the operands, the elliptic curve parameters must also be sent. This is 256 bits for the modulus of the curve, and 256 bits for the "a" parameter of the curve, which is needed for the point doubling operation. Once a curve is set, it will most likely be used many times, so this extra overhead of four instructions is also negligible.

One of the biggest challenges in hardware was making sure that the result of each operation did not

exceed the modulus. This means that on each addition becomes an add, a comparison, and potentially a subtraction. The modular multiplication then also consists of repeated addition and shifting instead of a large multiply followed by repeated subtraction of the modulus. The paths with the largest propagation delay have full adder cells, and it is likely due to the add-check-subtract operation with 256-bit fields, which is used numerous times throughout the functional units. This limits the clock period to 3.5 ns. It is not obvious how to best pipeline this due to the nature of the design, as there would be a tremendous cost in pipeline registers, but it should be explored in the future.

Currently, we only have a simulated version of the rocket interface working correctly. We have not been able to get the compilation of our interface with the rocket to correctly compile into Verilog code. It is unclear why the C++ backend works in this case while the Verilog does not.

Our hardware was written using Chisel. Verilog code was generated from Chisel and run through the Synopsys tool flow using a 32nm educational standard cell library. Multi-threshold cells were used, which cut the power numbers nearly in half.

## 7 Benchmarking Methodology

Currently, we only have a simulated version of the rocket interface working correctly. We have not been

able to get the compilation of our interface with the rocket to correctly compile into Verilog code. This means we couldn't actually run full signing tests through primetime to get power numbers. Instead we got ballpark power numbers by running the software stack through Spike with a simulator coprocessor to count the number of operations performed. We then measured the power of each operation (either Rocket integer operations or accelerator operations) with primetime, and summed up the total.

While this isn't the greatest result, it should give us reasonable accurate power numbers because of our algorithm. Essentially here we're relying on the fact that our algorithm never touches memory, which is what allows us to assume that the Rocket core will use a predictable amount of power.

The one big problem with this methodology is that it neglects the power used moving data between Rocket and the coprocessor – it just treats these cycles as a regular integer operation. Luckily, for any of the interesting configurations listed in the results section, this overhead should be minimal – easily within 0.1%. This is essentially because our software ECDSA implementation is so bad that the number of cycles run by Rocket whenever we are doing anything in software dwarfs the number of cycles moving data. Thus, even if moving data between Rocket and the coprocessor is very expensive relative to moving data within a Rocket core, it won't make much of a difference in the final power results.

## 8 Results

Platform	Power mJ/op	Speed op/sec	Area mm <sup>2</sup>
OpenSSL (45nm)	20	1000	
x86 (45 nm)	4000	5	
Rocket	800	0.05	
Virtex 2 (90nm)	4000	250	
Virtex 6 (45nm)	500	4000	
Mod Mul	200	0.3	0.04
Mod Mul+Div	2	20	0.10
Point Add+Db1	0.5	100	0.31
Point Mul	0.06	400	0.35

Table 2: Comparison of results to previous work

Table 2 shows a comparison of our work to previous FPGA implementations, as well as the design

space exploration of having different functionality in the hardware and software.

One important result to note is that certain blocks make sense to pair up in hardware. For example, a lot is gained from having modular multiplication and inversion in hardware as opposed to a single one. If the ECDSA algorithm is thought of as a hierarchy, then it can be seen that all layers in a hierarchy should be in hardware or none of them should be.

Another result is that it does it costs a significant amount to jump from the base of the hierarchy (multiply and invert), to the elliptic curve point and add layer. Add and double are complicated control blocks with many adders, and the area of the chip is tripled. However, to move from the add and double layer to the multiply layer, there is very little cost in area for a large improvement in energy/op. When the full point multiply is put in hardware, we start to approach the  $\mu\text{J}/\text{op}$  territory that would be desired for devices on the SwarmOS. With algorithmic and micro-architectural improvements, we will be even closer.

## 9 Future Work

Due to time constraints, there are many useful improvements that could be made to our accelerator. Elliptic field arithmetic is not an easy concept to grasp, and we had to make due with using only basic algorithms. One improvement that could be made on the elliptic field arithmetic is transforming a two-dimensional point into a three-dimensional representation for point multiplication. In this case, only one modular inversion is needed for the entire point multiplication (as opposed to hundreds). This further might make it unnecessary for an inversion hardware unit to be needed at all. Another improvement is to use Montgomery Multiplication[8]. This is a complex algorithm that performs successive multiplications fast after some initial overhead. If these algorithmic enhancements were made we would surely see a tremendous speedup. Both of these improvements have interesting time/space tradeoffs for hardware implementations, so could serve as additional design points for our accelerator.

A primary concern should be getting our software implementation up to snuff – this almost certainly means adopting OpenSSL as our testbench, hand-tuning its assembly for RISC-V, and then modifying it to be easily parameterizable so it can target our family of accelerators. We will then have to modify our hardware implementation to take into account

the algorithmic improvements that OpenSSL brings.

In addition to algorithmic improvements, there are a number of tricks one can play to squeeze performance out of ECDSA. The most important one appears to be to build hand-crafted routines for operating on particular curves – for example, a mod routine for NIST’s  $p256$  curve is hundreds of times faster than a generic one.

There are a number of architectural improvements that can be made to our coprocessor interface that can help mitigate the large overhead involved – either placing the accelerator closer to the pipeline or moving more state into the accelerator. Both of these lower the round-trip CPU to accelerator overhead and make some smaller accelerators possible, thus providing more interesting design points.

There are also a number of interesting micro-architectural improvements that can be made to our accelerator. We only have a 300MHz clock, which seems exceedingly slow for what should be a critical path of a single 256-bit adder. Almost certainly what happened was that there are actually two adders on the critical path, as there are cases where one step of an algorithm requires two additions back-to-back. Tweaking these state machines should provide an interesting trade-off between speed and power usage: essentially the higher clock rate will keep more transistors toggling, they just will not all be doing useful work.

Finally, a number of benchmarking improvements that would be interesting to see done. Specifically, it would be good to see a full-system benchmark running with our accelerator attached. This would be particularly interesting because it would allow for some degree of parallelism to be exposed by our accelerator to the benchmark. This could be very interesting because we may be able to provide light-weight accelerator contexts (compared to POSIX threads), which could result in an interesting userspace wrapper.

Another interesting aspect to explore would be an improved coprocessor interface. It’s possible that there could be a very light-weight accelerator that could provide significant energy-per-op improvements without requiring much overhead. Unfortunately we don’t think it will be possible to build this with the ROCC interface because of the inherent latency involved.

CS250 was a great course for learning how to design VLSI chips using the toolchain. It was difficult learning a new programming language to construct

hardware while also learning how to work with the VLSI toolflow. This could have been assisted with better tutorials on Chisel, specifically with tutorials on writing test harnesses for both the C++ simulator and more importantly the Verilog backend.

It would also have been helpful if instead of giving us ready-made directories and skeleton code in the early labs we were given step-by-step instructions on how to make a Chisel project from scratch and convert it to C++ and Verilog, because we mostly ended up hacking the the labs in really ugly ways with our accelerator to push it through the toolflow.

## References

- [1] Leonard Adleman. A subexponential algorithm for the discrete logarithm problem with applications to cryptography. 1979.
- [2] National Security Agency. The case for elliptic curve cryptography. 2009.
- [3] Taher ElGamal. A public key cryptosystem and a signature scheme based on discrete logarithms. *IEEE Transactions on Information Theory*, 1985.
- [4] Nadia Heninger, Zakir Durumeric, Eric Wustrow, and J. Alex Halderman. Mining your ps and qs: Detection of widespread weak keys in network devices, 2012.
- [5] Emilia Käsper. Fast elliptic curve cryptography in openssl.
- [6] Aqeel Khalique, Kuldeep Singh, and Sandeep Sood. Implementation of elliptic curve digital signature algorithm. *International Journal of Computer Applications*, 2010.
- [7] David W. Kravitz. Digital signature algorithm (us 5231668 a), 1991.
- [8] Ciaran J. McIvor, Marie McLoone, and John V. McCanny. Hardware elliptic curve cryptographic processor over  $gf(p)$ . *IEEE Transactions on Circuits and Systems*, 2006.
- [9] Ghanmy Nabil, Khelif Naziha, Fourati Lamia, and Kamoun Lotfi. Hardware implementation of elliptic curve digital signature algorithm (ecdsa) on koblitz curves. *8th IEEE, IET International Symposium on Communication Systems, Networks and Digital Signal Processing*, 2012.

- [10] National Institute of Standards and Technology (NIST). Digital signature standard (dss). *Federal Information Processing Standard, FIPS PUB 186-3*, 2009.